


Concept-flipping curves — attnlrp_epsilon · embed_dim · colored_mnist (n=500 imgs, N=384 detectors)
 $\Delta(n)$ = target-class prob after perturbing the n most/least relevant detectors, $\div$ baseline prob



— most-relevant-first: mean Δ over images
— least-relevant-first: mean Δ over images
 dark band: 95% bootstrap CI of the mean curve

faint band: ± 1 std across images (population spread)
● ● sign-crossing $n^\star$ on each curve (positive→negative relevance), matching colour
 $\Delta = 1$ (no change from baseline prediction)