


E2v2 — ViT-B/16 · ImageNet val: block- b activation outliers ending up in the final input-saliency mask, per method

A_b : $\|h_b(t)\|_2 > \mu_b + 4\sigma_b$ over the image's 196 patch tokens at the output of block b ; S_m : same $\mu + 4\sigma$ rule on the method's 196 per-patch saliency values



Annotation per bar: $c_{b,m}/n_b$ summed over $N=64$ correctly-classified images (prediction = label, so class-conditioning w.r.t. the predicted class). LRP = cp_lrp_baseline (CondAttribution, cond $\{y: [\hat{y}]\}$, sum $|R|$ per 16×16 patch); Chefer = grad-weighted attention rollout; attention rollout is class-agnostic by nature; occlusion = patch $\rightarrow$ image-mean, $\Delta p(\hat{y})$ clamped at 0. CLS excluded everywhere; h_b = forward hook on blocks $[b]$ output.